

exp13

November 9, 2025

1 Logistic Regression (from scratch)

1.1 1) Setup & Imports

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.metrics import accuracy_score, confusion_matrix,classification_report
```

1.2 2) Parameters

Change these and re-run the relevant cells below.

```
[2]: # ---- Parameters ----
DATA_PATH = 'loan_data.csv'      # Path to your CSV
TEST_SIZE = 0.2
RANDOM_STATE = 42

# Optimizer-specific defaults (you can override later or via widgets)
BGD_LEARNING_RATE = 0.1
BGD_EPOCHS = 1000

ADAM_LEARNING_RATE = 0.01
ADAM_EPOCHS = 1000
ADAM_BETA1 = 0.9
ADAM_BETA2 = 0.999
ADAM_EPSILON = 1e-8
```

1.3 3) From-Scratch Logistic Regression Class

```
[3]: # --- 1. The Logistic Regression Model Class ---  
  
class LogisticRegression:  
    """  
        A from-scratch implementation of Logistic Regression with two optimizers:  
        1. Batch Gradient Descent (BGD)  
        2. Adam Optimizer  
    """  
  
    def __init__(self, learning_rate=0.01, n_epochs=1000,  
                 beta1=0.9, beta2=0.999, epsilon=1e-8):  
        self.learning_rate = learning_rate  
        self.n_epochs = n_epochs  
        self.beta1 = beta1  
        self.beta2 = beta2  
        self.epsilon = epsilon  
        self.weights = None  
        self.bias = None  
        self.cost_history = []  
  
    def _sigmoid(self, z):  
        z_clipped = np.clip(z, -500, 500)  
        return 1 / (1 + np.exp(-z_clipped))  
  
    def _initialize_params(self, n_features):  
        self.weights = np.zeros(n_features)  
        self.bias = 0  
  
    def _compute_cost(self, y_true, y_pred):  
        m = y_true.shape[0]  
        epsilon = 1e-9  
        y_pred_clipped = np.clip(y_pred, epsilon, 1 - epsilon)  
        cost = -(1/m) * np.sum(  
            y_true * np.log(y_pred_clipped) + (1 - y_true) * np.log(1 - y_pred_clipped))  
        )  
        return cost  
  
    def fit_bgd(self, X, y):  
        print("Starting training with Batch Gradient Descent...")  
        m_samples, n_features = X.shape  
        self._initialize_params(n_features)  
        self.cost_history = []  
  
        for i in range(self.n_epochs):
```

```

linear_model = X @ self.weights + self.bias
y_pred = self._sigmoid(linear_model)
cost = self._compute_cost(y, y_pred)
self.cost_history.append(cost)
dw = (1 / m_samples) * (X.T @ (y_pred - y))
db = (1 / m_samples) * np.sum(y_pred - y)
self.weights -= self.learning_rate * dw
self.bias -= self.learning_rate * db
if (i+1) % 100 == 0:
    print(f"BGD Epoch {i+1}/{self.n_epochs}, Cost: {cost:.4f}")

def fit_adam(self, X, y):
    print("Starting training with Adam Optimizer...")
    m_samples, n_features = X.shape
    self._initialize_params(n_features)
    self.cost_history = []
    m_w, v_w, m_b, v_b = 0, 0, 0, 0

    for i in range(self.n_epochs):
        t = i + 1
        linear_model = X @ self.weights + self.bias
        y_pred = self._sigmoid(linear_model)
        cost = self._compute_cost(y, y_pred)
        self.cost_history.append(cost)
        dw = (1 / m_samples) * (X.T @ (y_pred - y))
        db = (1 / m_samples) * np.sum(y_pred - y)
        m_w = self.beta1 * m_w + (1 - self.beta1) * dw
        m_b = self.beta1 * m_b + (1 - self.beta1) * db
        v_w = self.beta2 * v_w + (1 - self.beta2) * (dw**2)
        v_b = self.beta2 * v_b + (1 - self.beta2) * (db**2)
        m_w_corr = m_w / (1 - self.beta1**t)
        m_b_corr = m_b / (1 - self.beta1**t)
        v_w_corr = v_w / (1 - self.beta2**t)
        v_b_corr = v_b / (1 - self.beta2**t)
        self.weights -= self.learning_rate * m_w_corr / (np.sqrt(v_w_corr) +
+ self.epsilon)
        self.bias -= self.learning_rate * m_b_corr / (np.sqrt(v_b_corr) +
+ self.epsilon)
        if (i+1) % 100 == 0:
            print(f"Adam Epoch {i+1}/{self.n_epochs}, Cost: {cost:.4f}")

    def predict_proba(self, X):
        linear_model = X @ self.weights + self.bias
        return self._sigmoid(linear_model)

    def predict(self, X, threshold=0.5):
        probas = self.predict_proba(X)

```

```
    return (probas >= threshold).astype(int)
```

1.4 4) Helper Functions

```
[4]: def plot_confusion_matrix(y_true, y_pred, title):
    cm = confusion_matrix(y_true, y_pred)
    plt.figure(figsize=(6, 4))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
                xticklabels=['Rejected (0)', 'Approved (1)'],
                yticklabels=['Rejected (0)', 'Approved (1)'])
    plt.xlabel('Predicted')
    plt.ylabel('Actual')
    plt.title(title)
    plt.show()

def evaluate_and_report(model_name, y_true, y_pred):
    acc = accuracy_score(y_true, y_pred)
    print(f"\n{model_name} Accuracy: {acc:.4f}")
    print(f"{model_name} Classification Report:")
    print(classification_report(y_true, y_pred, target_names=['Rejected (0)', 'Approved (1)']))
    return acc
```

1.5 5) Load Data

```
[6]: # Make sure DATA_PATH points to your CSV.
try:
    data = pd.read_csv(r"C:\\\\Sarvesh\\\\Mtech\\\\Foundation of Data\\\\
    ↪Engineering\\\\Assignments\\\\EX 13\\\\loan_data.csv")
except FileNotFoundError:
    raise FileNotFoundError("Error: '{r'C:\\\\Sarvesh\\\\Mtech\\\\Foundation of Data\\\\
    ↪Engineering\\\\Assignments\\\\EX 13\\\\loan_data.csv'}' not found. Place it next\\\\
    ↪to this notebook or update DATA_PATH.")

print('Data loaded successfully. Shape:', data.shape)
# Drop rows with missing target
data.dropna(subset=['loan_status'], inplace=True)
X = data.drop('loan_status', axis=1)
y = data['loan_status'].values

print('Preview features:')
display(X.head(3))
print('Preview target:')
display(pd.Series(y).head(3))
```

Data loaded successfully. Shape: (45000, 14)

Preview features:

```

person_age person_gender person_education person_income person_emp_exp \
0      22.0      female        Master     71948.0          0
1      21.0      female    High School   12282.0          0
2      25.0      female    High School   12438.0          3

person_home_ownership loan_amnt loan_intent loan_int_rate \
0           RENT    35000.0    PERSONAL     16.02
1           OWN     1000.0    EDUCATION    11.14
2      MORTGAGE    5500.0    MEDICAL     12.87

loan_percent_income cb_person_cred_hist_length credit_score \
0            0.49             3.0       561
1            0.08             2.0       504
2            0.44             3.0       635

previous_loan_defaults_on_file
0                  No
1                 Yes
2                  No

Preview target:
0    1
1    0
2    1
dtype: int64

```

1.6 6) Build Preprocessing Pipelines

```
[7]: numeric_features = X.select_dtypes(include=np.number).columns
categorical_features = X.select_dtypes(include='object').columns
print(f"Numerical features: {list(numeric_features)}")
print(f"Categorical features: {list(categorical_features)}")

numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())
])
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('onehot', OneHotEncoder(handle_unknown='ignore'))
])

preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)
    ],
)
```

```

        remainder='passthrough'
    )

Numerical features: ['person_age', 'person_income', 'person_emp_exp',
'loan_amnt', 'loan_int_rate', 'loan_percent_income',
'cb_person_cred_hist_length', 'credit_score']
Categorical features: ['person_gender', 'person_education',
'person_home_ownership', 'loan_intent', 'previous_loan_defaults_on_file']

```

1.7 7) Train/Test Split & Transform

```
[8]: X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=TEST_SIZE, random_state=RANDOM_STATE, stratify=y
)
print(f"Training set size: {X_train.shape[0]}")
print(f"Test set size: {X_test.shape[0]}")

print('Fitting preprocessor and transforming training data...')
X_train_processed = preprocessor.fit_transform(X_train)

print('Transforming test data...')
X_test_processed = preprocessor.transform(X_test)

if not isinstance(X_train_processed, np.ndarray):
    X_train_processed = X_train_processed.toarray()
    X_test_processed = X_test_processed.toarray()

print('Preprocessing complete. Shapes:', X_train_processed.shape, ▶
    X_test_processed.shape)
```

```

Training set size: 36000
Test set size: 9000
Fitting preprocessor and transforming training data...
Transforming test data...
Preprocessing complete. Shapes: (36000, 27) (9000, 27)
```

1.8 8) Train — Batch Gradient Descent (BGD)

```
[9]: model_bgd = LogisticRegression(learning_rate=BGD_LEARNING_RATE, ▶
    n_epochs=BGD_EPOCHS)
model_bgd.fit_bgd(X_train_processed, y_train)
```

```

Starting training with Batch Gradient Descent...
BGD Epoch 100/1000, Cost: 0.3134
BGD Epoch 200/1000, Cost: 0.2761
BGD Epoch 300/1000, Cost: 0.2605
BGD Epoch 400/1000, Cost: 0.2518
BGD Epoch 500/1000, Cost: 0.2462
```

```
BGD Epoch 600/1000, Cost: 0.2423  
BGD Epoch 700/1000, Cost: 0.2394  
BGD Epoch 800/1000, Cost: 0.2372  
BGD Epoch 900/1000, Cost: 0.2354  
BGD Epoch 1000/1000, Cost: 0.2340
```


1.9 9) Evaluate — BGD

```
[10]: y_pred_bgd = model_bgd.predict(X_test_processed)  
acc_bgd = evaluate_and_report('BGD Model', y_test, y_pred_bgd)  
plot_confusion_matrix(y_test, y_pred_bgd, 'BGD Confusion Matrix')
```

BGD Model Accuracy: 0.8982

BGD Model Classification Report:

	precision	recall	f1-score	support
Rejected (0)	0.93	0.94	0.94	7000
Approved (1)	0.79	0.74	0.76	2000
accuracy			0.90	9000
macro avg	0.86	0.84	0.85	9000
weighted avg	0.90	0.90	0.90	9000



1.10 10) Train — Adam Optimizer

```
[11]: model_adam = LogisticRegression(  
    learning_rate=ADAM_LEARNING_RATE,  
    n_epochs=ADAM_EPOCHS,  
    beta1=ADAM_BETA1, beta2=ADAM_BETA2, epsilon=ADAM_EPSILON  
)  
model_adam.fit_adam(X_train_processed, y_train)
```

Starting training with Adam Optimizer...

Adam Epoch 100/1000, Cost: 0.2927
Adam Epoch 200/1000, Cost: 0.2544
Adam Epoch 300/1000, Cost: 0.2402
Adam Epoch 400/1000, Cost: 0.2332
Adam Epoch 500/1000, Cost: 0.2295
Adam Epoch 600/1000, Cost: 0.2274
Adam Epoch 700/1000, Cost: 0.2261
Adam Epoch 800/1000, Cost: 0.2252
Adam Epoch 900/1000, Cost: 0.2246
Adam Epoch 1000/1000, Cost: 0.2241


1.11 11) Evaluate — Adam

```
[12]: y_pred_adam = model_adam.predict(X_test_processed)  
acc_adam = evaluate_and_report('Adam Model', y_test, y_pred_adam)  
plot_confusion_matrix(y_test, y_pred_adam, 'Adam Optimizer Confusion Matrix')
```

Adam Model Accuracy: 0.8989

Adam Model Classification Report:

	precision	recall	f1-score	support
Rejected (0)	0.93	0.94	0.94	7000
Approved (1)	0.79	0.74	0.77	2000
accuracy			0.90	9000
macro avg	0.86	0.84	0.85	9000
weighted avg	0.90	0.90	0.90	9000



1.12 12) Compare Cost Histories

```
[13]: plt.figure(figsize=(10,5))
plt.plot(model_bgd.cost_history, label=f'BGD (LR={BGD_LEARNING_RATE})')
plt.plot(model_adam.cost_history, label=f'Adam (LR={ADAM_LEARNING_RATE})',  
         linestyle='--')
plt.title('Cost Function Over Epochs')
plt.xlabel('Epoch')
plt.ylabel('Log Loss (Cost)')
plt.legend()
plt.ylim(0, 1)
plt.grid(True)
plt.show()
```

